

Symptom_Selection

April 26, 2025

```
[ ]: ! pip install seaborn
```

```
Requirement already satisfied: seaborn in
/Users/zhengfeibian/anaconda3/envs/happy/lib/python3.11/site-packages (0.13.2)
Requirement already satisfied: numpy!=1.24.0,>=1.20 in
/Users/zhengfeibian/anaconda3/envs/happy/lib/python3.11/site-packages (from
seaborn) (2.2.5)
Requirement already satisfied: pandas>=1.2 in
/Users/zhengfeibian/anaconda3/envs/happy/lib/python3.11/site-packages (from
seaborn) (2.2.3)
Requirement already satisfied: matplotlib!=3.6.1,>=3.4 in
/Users/zhengfeibian/anaconda3/envs/happy/lib/python3.11/site-packages (from
seaborn) (3.10.1)
Requirement already satisfied: contourpy>=1.0.1 in
/Users/zhengfeibian/anaconda3/envs/happy/lib/python3.11/site-packages (from
matplotlib!=3.6.1,>=3.4->seaborn) (1.3.2)
Requirement already satisfied: cycler>=0.10 in
/Users/zhengfeibian/anaconda3/envs/happy/lib/python3.11/site-packages (from
matplotlib!=3.6.1,>=3.4->seaborn) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in
/Users/zhengfeibian/anaconda3/envs/happy/lib/python3.11/site-packages (from
matplotlib!=3.6.1,>=3.4->seaborn) (4.57.0)
Requirement already satisfied: kiwisolver>=1.3.1 in
/Users/zhengfeibian/anaconda3/envs/happy/lib/python3.11/site-packages (from
matplotlib!=3.6.1,>=3.4->seaborn) (1.4.8)
Requirement already satisfied: packaging>=20.0 in
/Users/zhengfeibian/anaconda3/envs/happy/lib/python3.11/site-packages (from
matplotlib!=3.6.1,>=3.4->seaborn) (24.2)
Requirement already satisfied: pillow>=8 in
/Users/zhengfeibian/anaconda3/envs/happy/lib/python3.11/site-packages (from
matplotlib!=3.6.1,>=3.4->seaborn) (11.2.1)
Requirement already satisfied: pyparsing>=2.3.1 in
/Users/zhengfeibian/anaconda3/envs/happy/lib/python3.11/site-packages (from
matplotlib!=3.6.1,>=3.4->seaborn) (3.2.3)
Requirement already satisfied: python-dateutil>=2.7 in
/Users/zhengfeibian/anaconda3/envs/happy/lib/python3.11/site-packages (from
matplotlib!=3.6.1,>=3.4->seaborn) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in
```

```

/Users/zhengfeibian/anaconda3/envs/happy/lib/python3.11/site-packages (from
pandas>=1.2->seaborn) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in
/Users/zhengfeibian/anaconda3/envs/happy/lib/python3.11/site-packages (from
pandas>=1.2->seaborn) (2025.2)
Requirement already satisfied: six>=1.5 in
/Users/zhengfeibian/anaconda3/envs/happy/lib/python3.11/site-packages (from
python-dateutil>=2.7->matplotlib!=3.6.1,>=3.4->seaborn) (1.17.0)

```

```

[ ]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix, \
    roc_auc_score, accuracy_score
import matplotlib.pyplot as plt
import seaborn as sns
import os

```

```

[ ]: ## Prepare the symptom prediction dataset

# Load the Alzheimer's dataset
# Dataset source: https://www.kaggle.com/datasets/rabieelkharoua/
    ↪alzhimers-disease-dataset
data_path = "/Users/zhengfeibian/Desktop/5630final/MyOwnChooseDataSets/
    ↪alzhimers_disease_data.csv"

df = pd.read_csv(data_path)

# Preview the first few rows
df.head()

```

```

[ ]:

```

	PatientID	Age	Gender	Ethnicity	EducationLevel	BMI	Smoking	\
0	4751	73	0	0	2	22.927749	0	
1	4752	89	0	0	0	26.827681	0	
2	4753	73	0	3	1	17.795882	0	
3	4754	74	1	0	1	33.800817	1	
4	4755	89	0	0	0	20.716974	0	

	AlcoholConsumption	PhysicalActivity	DietQuality	...	MemoryComplaints	\
0	13.297218	6.327112	1.347214	...	0	
1	4.542524	7.619885	0.518767	...	0	
2	19.555085	7.844988	1.826335	...	0	
3	12.209266	8.428001	7.435604	...	0	
4	18.454356	6.310461	0.795498	...	0	

	BehavioralProblems	ADL	Confusion	Disorientation	\
0	0	1.725883	0	0	
1	0	2.592424	0	0	
2	0	7.119548	0	1	
3	1	6.481226	0	0	
4	0	0.014691	0	0	

	PersonalityChanges	DifficultyCompletingTasks	Forgetfulness	Diagnosis	\
0	0		1	0	0
1	0		0	1	0
2	0		1	0	0
3	0		0	0	0
4	1		1	0	0

	DoctorInCharge
0	XXXConfid
1	XXXConfid
2	XXXConfid
3	XXXConfid
4	XXXConfid

[5 rows x 35 columns]

```
[ ]: # Print all column names
print(list(df.columns))
```

```
['PatientID', 'Age', 'Gender', 'Ethnicity', 'EducationLevel', 'BMI', 'Smoking',
'AlcoholConsumption', 'PhysicalActivity', 'DietQuality', 'SleepQuality',
'FamilyHistoryAlzheimers', 'CardiovascularDisease', 'Diabetes', 'Depression',
'HeadInjury', 'Hypertension', 'SystolicBP', 'DiastolicBP', 'CholesterolTotal',
'CholesterolLDL', 'CholesterolHDL', 'CholesterolTriglycerides', 'MMSE',
'FunctionalAssessment', 'MemoryComplaints', 'BehavioralProblems', 'ADL',
'Confusion', 'Disorientation', 'PersonalityChanges',
'DifficultyCompletingTasks', 'Forgetfulness', 'Diagnosis', 'DoctorInCharge']
```

```
[ ]: # Data Cleaning
# Drop irrelevant columns like Patient ID and Doctor in Charge
df_clean = df.drop(columns=["PatientID", "DoctorInCharge"])

# Define Target Symptoms
# We will predict multiple symptoms separately
target_variables = ["Forgetfulness", "Confusion", "Disorientation",
↪ "PersonalityChanges"]

# Set Output Directory
output_dir = "/Users/zhengfeibian/Desktop/5630final/MyOwnChooseDataSets"
```

```

# For each target variable, prepare training and testing sets
for target_variable in target_variables:
    print(f"\nPreparing data for target variable: {target_variable}")

    # Split features (X) and label (y)
    X = df_clean.drop(columns=target_variables) # Features (drop all targets)
    y = df_clean[target_variable] # Current label

    # Split into train and test sets (80/20) with stratification
    X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=0.2, random_state=42, stratify=y
    )

    # Merge features and labels
    train_df = X_train.copy()
    train_df[target_variable] = y_train

    test_df = X_test.copy()
    test_df[target_variable] = y_test

    # Save train and test sets
    train_path = os.path.join(output_dir, f"train_{target_variable.lower()}.
    ↪csv")
    test_path = os.path.join(output_dir, f"test_{target_variable.lower()}.csv")

    train_df.to_csv(train_path, index=False)
    test_df.to_csv(test_path, index=False)

    print(f"Saved train set: {train_path}")
    print(f"Saved test set: {test_path}")

```

Preparing data for target variable: Forgetfulness

Saved train set: /Users/zhengfeibian/Desktop/5630final/MyOwnChooseDataSets/train_forgetfulness.csv

Saved test set:

/Users/zhengfeibian/Desktop/5630final/MyOwnChooseDataSets/test_forgetfulness.csv

Preparing data for target variable: Confusion

Saved train set:

/Users/zhengfeibian/Desktop/5630final/MyOwnChooseDataSets/train_confusion.csv

Saved test set:

/Users/zhengfeibian/Desktop/5630final/MyOwnChooseDataSets/test_confusion.csv

Preparing data for target variable: Disorientation

Saved train set: /Users/zhengfeibian/Desktop/5630final/MyOwnChooseDataSets/train_disorientation.csv

Saved test set: /Users/zhengfeibian/Desktop/5630final/MyOwnChooseDataSets/test_d

isorientation.csv

Preparing data for target variable: PersonalityChanges

Saved train set: /Users/zhengfeibian/Desktop/5630final/MyOwnChooseDataSets/train_personalitychanges.csv

Saved test set: /Users/zhengfeibian/Desktop/5630final/MyOwnChooseDataSets/test_personalitychanges.csv

```
[ ]: # Symptom occurrence statistics
# Count number of patients showing each symptom
symptom_counts = df_clean[["Forgetfulness", "Confusion", "Disorientation",
↪ "PersonalityChanges"]].sum()
print(symptom_counts)
```

```
Forgetfulness      648
Confusion          441
Disorientation     340
PersonalityChanges 324
dtype: int64
```

```
[ ]: from sklearn.preprocessing import StandardScaler
```

```
[ ]: # Train Logistic Regression Models
output_dir = "/Users/zhengfeibian/Desktop/5630final/MyOwnChooseDataSets"

# To store logistic regression results
results = []

for target_variable in target_variables:
    print(f"\nTraining Logistic Regression for target: {target_variable}")

    # Load train and test sets
    train_path = f"{output_dir}/train_{target_variable.lower()}.csv"
    test_path = f"{output_dir}/test_{target_variable.lower()}.csv"

    train_df = pd.read_csv(train_path)
    test_df = pd.read_csv(test_path)

    # Separate features and label
    X_train = train_df.drop(columns=[target_variable])
    y_train = train_df[target_variable]
    X_test = test_df.drop(columns=[target_variable])
    y_test = test_df[target_variable]

    # Standardize features
    scaler = StandardScaler()
    X_train_scaled = scaler.fit_transform(X_train)
    X_test_scaled = scaler.transform(X_test)
```

```

# Build and train Logistic Regression model
lr_model = LogisticRegression(class_weight='balanced', max_iter=1000,
random_state=42)
lr_model.fit(X_train_scaled, y_train)

# Predict
y_pred = lr_model.predict(X_test_scaled)
y_prob = lr_model.predict_proba(X_test_scaled)[: , 1]

# Evaluate
acc = accuracy_score(y_test, y_pred)
auc = roc_auc_score(y_test, y_prob)
report = classification_report(y_test, y_pred, digits=4)
cm = confusion_matrix(y_test, y_pred)

print(f"Accuracy: {acc:.4f}")
print(f"ROC-AUC: {auc:.4f}")
print(f"Classification Report:\n{report}")
print(f"Confusion Matrix:\n{cm}")

# Save performance
results.append({
    "Target": target_variable,
    "Accuracy": acc,
    "ROC_AUC": auc
})

```

Training Logistic Regression for target: Forgetfulness

Accuracy: 0.4977

ROC-AUC: 0.4721

Classification Report:

	precision	recall	f1-score	support
0	0.6736	0.5433	0.6015	300
1	0.2713	0.3923	0.3208	130
accuracy			0.4977	430
macro avg	0.4724	0.4678	0.4611	430
weighted avg	0.5519	0.4977	0.5166	430

Confusion Matrix:

```
[[163 137]
```

```
[ 79  51]]
```

Training Logistic Regression for target: Confusion

Accuracy: 0.4628

ROC-AUC: 0.4671

Classification Report:

	precision	recall	f1-score	support
0	0.7789	0.4532	0.5730	342
1	0.1905	0.5000	0.2759	88
accuracy			0.4628	430
macro avg	0.4847	0.4766	0.4244	430
weighted avg	0.6585	0.4628	0.5122	430

Confusion Matrix:

```
[[155 187]
 [ 44  44]]
```

Training Logistic Regression for target: Disorientation

Accuracy: 0.5326

ROC-AUC: 0.5375

Classification Report:

	precision	recall	f1-score	support
0	0.8676	0.5249	0.6540	362
1	0.1848	0.5735	0.2796	68
accuracy			0.5326	430
macro avg	0.5262	0.5492	0.4668	430
weighted avg	0.7596	0.5326	0.5948	430

Confusion Matrix:

```
[[190 172]
 [ 29  39]]
```

Training Logistic Regression for target: PersonalityChanges

Accuracy: 0.5349

ROC-AUC: 0.4395

Classification Report:

	precision	recall	f1-score	support
0	0.8367	0.5616	0.6721	365
1	0.1351	0.3846	0.2000	65
accuracy			0.5349	430
macro avg	0.4859	0.4731	0.4361	430
weighted avg	0.7307	0.5349	0.6008	430

Confusion Matrix:

```
[[205 160]
 [ 40  25]]
```

```
[ ]: # Summary of Logistic Regression Results
import pandas as pd

lr_results_df = pd.DataFrame(results)
print("\nLogistic Regression Results Summary:")
print(lr_results_df)
```

Logistic Regression Results Summary:

	Target	Accuracy	ROC_AUC
0	Forgetfulness	0.497674	0.472051
1	Confusion	0.462791	0.467072
2	Disorientation	0.532558	0.537496
3	PersonalityChanges	0.534884	0.439452

```
[ ]: # Train Random Forest Models
output_dir = "/Users/zhengfeibian/Desktop/5630final/MyOwnChooseDataSets"

# To store random forest results
rf_results = []

for target_variable in target_variables:
    print(f"\nTraining Random Forest for target: {target_variable}")

    # Load train and test sets
    train_path = f"{output_dir}/train_{target_variable.lower()}.csv"
    test_path = f"{output_dir}/test_{target_variable.lower()}.csv"

    train_df = pd.read_csv(train_path)
    test_df = pd.read_csv(test_path)

    # Separate features and label
    X_train = train_df.drop(columns=[target_variable])
    y_train = train_df[target_variable]
    X_test = test_df.drop(columns=[target_variable])
    y_test = test_df[target_variable]

    # Random Forest does not require feature scaling
    rf_model = RandomForestClassifier(
        n_estimators=100,
        max_depth=None,
        class_weight='balanced',
        random_state=42
    )
    rf_model.fit(X_train, y_train)

    # Predict
    y_pred = rf_model.predict(X_test)
```



```

y_prob = rf_model.predict_proba(X_test)[: , 1]

# Evaluate
acc = accuracy_score(y_test, y_pred)
auc = roc_auc_score(y_test, y_prob)
report = classification_report(y_test, y_pred, digits=4)
cm = confusion_matrix(y_test, y_pred)

print(f"Accuracy: {acc:.4f}")
print(f"ROC-AUC: {auc:.4f}")
print(f"Classification Report:\n{report}")
print(f"Confusion Matrix:\n{cm}")

# Save performance
rf_results.append({
    "Target": target_variable,
    "Accuracy": acc,
    "ROC_AUC": auc
})

```

Training Random Forest for target: Forgetfulness

Accuracy: 0.6977

ROC-AUC: 0.5057

Classification Report:

	precision	recall	f1-score	support
0	0.6986	0.9967	0.8214	300
1	0.5000	0.0077	0.0152	130
accuracy			0.6977	430
macro avg	0.5993	0.5022	0.4183	430
weighted avg	0.6386	0.6977	0.5777	430

Confusion Matrix:

```

[[299  1]
 [129  1]]

```

Training Random Forest for target: Confusion

/Users/zhengfeibian/anaconda3/envs/happy/lib/python3.11/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```

_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/Users/zhengfeibian/anaconda3/envs/happy/lib/python3.11/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted

```

samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

/Users/zhengfeibian/anaconda3/envs/happy/lib/python3.11/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

Accuracy: 0.7953

ROC-AUC: 0.4596

Classification Report:

	precision	recall	f1-score	support
0	0.7953	1.0000	0.8860	342
1	0.0000	0.0000	0.0000	88
accuracy			0.7953	430
macro avg	0.3977	0.5000	0.4430	430
weighted avg	0.6326	0.7953	0.7047	430

Confusion Matrix:

```
[[342  0]
 [ 88  0]]
```

Training Random Forest for target: Disorientation

```
/Users/zhengfeibian/anaconda3/envs/happy/lib/python3.11/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

```
/Users/zhengfeibian/anaconda3/envs/happy/lib/python3.11/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

```
/Users/zhengfeibian/anaconda3/envs/happy/lib/python3.11/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

Accuracy: 0.8419

ROC-AUC: 0.4852

Classification Report:

	precision	recall	f1-score	support
0	0.8419	1.0000	0.9141	362
1	0.0000	0.0000	0.0000	68

accuracy			0.8419	430
macro avg	0.4209	0.5000	0.4571	430
weighted avg	0.7087	0.8419	0.7696	430

Confusion Matrix:

```
[[362  0]
 [ 68  0]]
```

Training Random Forest for target: PersonalityChanges

Accuracy: 0.8488

ROC-AUC: 0.5329

Classification Report:

	precision	recall	f1-score	support
0	0.8488	1.0000	0.9182	365
1	0.0000	0.0000	0.0000	65

accuracy			0.8488	430
macro avg	0.4244	0.5000	0.4591	430
weighted avg	0.7205	0.8488	0.7794	430

Confusion Matrix:

```
[[365  0]
 [ 65  0]]
```

/Users/zhengfeibian/anaconda3/envs/happy/lib/python3.11/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

/Users/zhengfeibian/anaconda3/envs/happy/lib/python3.11/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

/Users/zhengfeibian/anaconda3/envs/happy/lib/python3.11/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

Among the Alzheimer's disease symptoms analyzed, Forgetfulness shows the highest prevalence, suggesting it should be prioritized when designing keyword extraction strategies. Disorientation and Personality Changes, although less prevalent, demonstrated higher predictability, indicating their clear pattern within patient characteristics."

```
[ ]: # Summary of Random Forest Results
import pandas as pd

rf_results_df = pd.DataFrame(rf_results)
print("\nRandom Forest Results Summary:")
print(rf_results_df)
```

```
Random Forest Results Summary:
      Target  Accuracy  ROC_AUC
0  Forgetfulness  0.697674  0.505731
1      Confusion  0.795349  0.459629
2  Disorientation  0.841860  0.485172
3 PersonalityChanges  0.848837  0.532919
```